

AI Parameters Scoring

1. Purpose and Overview

This document defines the scoring methodology used to evaluate and prioritise AI use cases within Oracle client engagements. Each use case is assessed across six dimensions. Five dimensions contribute positive scores; one dimension (Complexity) applies a penalty. The resulting weighted score drives prioritisation decisions.

The framework is not a generic AI evaluation tool. It is an Oracle Advisory filtering framework. If a use case does not fit the Oracle ERP ecosystem, it is not relevant, regardless of its standalone AI merit. Oracle Fit is therefore the primary gating factor.

Dimension	What it measures	Weight	Role
Oracle fit	Alignment with Oracle Cloud ecosystem — data, services, integration, pipeline	35%	Primary gate
Business value	Measurable impact — revenue, cost, risk, efficiency, strategic advantage	25%	ROI anchor
Data readiness	How prepared, accessible and usable the data is for building and deploying AI	15%	Feasibility
Complexity	How difficult it is to implement end-to-end in an enterprise environment	-10%	Penalty
Client demand	How strongly and frequently enterprises are requesting or adopting this use case	7.5%	Market signal
Time to value	How quickly measurable ROI can be delivered after implementation begins	7.5%	Speed signal

Weighted score formula:

$$\text{Score} = (\text{Oracle fit} \times 0.35) + (\text{Business value} \times 0.25) + (\text{Data readiness} \times 0.15) - (\text{Complexity} \times 0.10) + (\text{Client demand} \times 0.075) + (\text{Time to value} \times 0.075)$$

2. Oracle Fit

Oracle fit measures how well a use case can be implemented using the Oracle Cloud ecosystem with minimal gaps and maximum advantage. It is not a binary assessment. It is a spectrum from complete Oracle-native delivery to cases where Oracle adds negligible value.

Weight: 35% (primary gate). If Oracle fit is low, the use case is out of scope for an Oracle Advisory engagement regardless of how high other scores are.

Oracle Cloud services in scope

ADW (Autonomous Data Warehouse) · OCI Data Flow · OCI Data Science · OCI AI Services · OCI Observability · OCI DevOps (deployment stage)

Scoring dimensions

Dimension	What is assessed
1. Data residency	Is data already within Oracle systems (ERP, HCM, SCM, Oracle DB)? Data originating from Oracle is immediately consumable without external extraction pipelines.
2. Native service coverage	Can the use case be implemented using ADW, OCI Data Flow, OCI Data Science, and OCI AI Services without external tools?
3. Integration strength	How easily can the model connect with Oracle ERP, SCM, HCM, and existing workflows? Strong integration avoids custom middleware and reduces risk.
4. End-to-end capability	Can OCI handle the complete pipeline: Data → Model → Deployment → Monitoring? Partial capability requires external tools and increases delivery complexity.
5. Differentiation advantage	Does Oracle offer a meaningful competitive advantage versus AWS, Azure, or GCP for this specific use case? Strong differentiation justifies the Oracle-first advisory approach.

Scale (0–10)

Score	Level	Characteristics
10	Perfect fit	Data fully in Oracle. All services native. Seamless ERP/SCM/HCM integration. Complete OCI pipeline. Strong differentiation. Example: Invoice automation, ERP-based AI.
9	Very strong fit	Mostly Oracle data. Almost full service coverage. Strong integration. Minor pipeline gaps. High differentiation advantage.
8	Strong fit	Partial Oracle data. Mostly covered. Good integration. Small external dependency. Moderate–high advantage.
7	Moderately strong	Mixed data sources. Core services covered. Moderate integration. Needs some external tools. Moderate advantage.
6	Medium fit	Mostly external data. Partial coverage. Limited integration. Additional tools required. Neutral advantage.
5	Moderate fit	External data. Limited OCI usage. Minimal integration. Fragmented pipeline. Low advantage.
4	Low–moderate	Data not in Oracle. Only 1–2 services usable. Weak integration. Mostly external stack. Very low advantage.
3	Low fit	External ecosystem. Minimal service relevance. Poor integration. Not feasible in OCI alone. No advantage.

2	Very low fit	Completely outside Oracle. Services barely usable. Integration not meaningful. External stack required.
1	Negligible fit	No Oracle relevance. Services not applicable. OCI delivery impossible. No advantage.
0	No fit	Completely unrelated to Oracle ecosystem. Do not proceed.

Worked examples

Factor	Invoice anomaly detection	Customer chatbot
Data	Fully in Oracle ERP	External systems
Services	Full OCI coverage	Limited OCI relevance
Integration	Strong / seamless	Weak
Oracle fit score	10	6–7

3. Business Value

Business value measures the measurable impact the AI use case creates for the organisation. It encompasses revenue increase, cost reduction, risk reduction, efficiency gains, and strategic advantage.

Weight: 25%. Advisory is ROI-driven. Clients care about revenue and cost. High value within the Oracle ecosystem is required, not just high technical merit.

Scoring dimensions

Dimension	What is assessed
1. Financial impact	Does this use case directly increase revenue, reduce costs, or reduce risk? Magnitude and directness of the financial linkage.
2. Scope of impact	Is the benefit confined to a small team, a single department, one major function, multiple business units, or organisation-wide?
3. Process criticality	Does the use case improve a core business process (directly revenue or cost-generating) or a support function?
4. Frequency of use	How often is the AI output consumed — daily and continuous, frequent, regular, occasional, or rare?

Scale (0–10)

Score	Level	Combination
10	Transformational	Financial: Very high (direct revenue ↑ or major cost ↓). Scope: Organisation-wide. Criticality: Core business function. Frequency: Continuous / daily.
9	Very high value	Financial: Very high. Scope: Multiple business units. Criticality: Core function. Frequency: Frequent.
8	High value	Financial: High. Scope: One major function. Criticality: Core or near-core. Frequency: Frequent.
7	Moderately high	Financial: Moderate–high. Scope: One function. Criticality: Important but not core. Frequency: Frequent.
6	Medium value	Financial: Moderate (efficiency gains). Scope: One team / department. Criticality: Support function. Frequency: Regular.
5	Moderate value	Financial: Low–moderate. Scope: Limited team. Criticality: Support process. Frequency: Occasional.
4	Low–moderate	Financial: Low. Scope: Small team. Criticality: Non-critical. Frequency: Occasional.
3	Low value	Financial: Minimal. Scope: Individual / team level. Criticality: Non-critical. Frequency: Rare.
2	Very low value	Financial: Negligible. Scope: Very limited. Criticality: Optional process. Frequency: Rare.
1	Negligible value	Financial: Almost none. Scope: Individual use. Criticality: Not important. Frequency: Very rare.
0	No value	No measurable business impact.

4. Data Readiness

Data readiness measures how prepared, accessible, and usable the data is for building and deploying an AI solution. This is not simply whether data exists. It encompasses availability, quality, structure, accessibility, and governance.

Weight: 15%. The strongest insight from this project is that AI fails due to data. Even if Oracle fit is high, poor data readiness leads to failure. Data readiness is the most common failure point in enterprise AI delivery.

Scoring dimensions

Dimension	What is assessed
1. Data availability	Is sufficient historical and current data present, or does it need to be collected from scratch?
2. Data quality	Is the data clean, complete, and consistent, or does it require significant preparation?
3. Data structure	Is data structured (SQL tables, CSVs) or unstructured (PDFs, emails, images, text)?
4. Data accessibility	Is data centralised and easily accessible, or siloed across systems with restricted access?
5. Data governance	Are ownership, security, and compliance controls defined and operational?

Scale (0–10)

Score	Level	Characteristics
10	Fully ready	Availability: Complete historical data. Quality: Clean, minimal missing values. Structure: Fully structured. Accessibility: Centralised. Governance: Strong.
9	Very high	Availability: Mostly complete. Quality: Minor issues. Structure: Mostly structured. Accessibility: Good access. Governance: Well defined.
8	High readiness	Availability: Sufficient. Quality: Some cleaning required. Structure: Structured + small unstructured. Accessibility: Accessible with effort. Governance: Moderate.
7	Moderately high	Availability: Available but limited history. Quality: Moderate inconsistencies. Structure: Mixed. Accessibility: Some silos. Governance: Partial.
6	Medium readiness	Availability: Partial. Quality: Noticeable issues. Structure: Mixed / semi-structured. Accessibility: Limited. Governance: Weak.
5	Moderate readiness	Availability: Incomplete. Quality: Significant cleaning required. Structure: Mostly unstructured. Accessibility: Siloed. Governance: Minimal.
4	Low–moderate	Availability: Data scattered. Quality: Poor. Structure: Unstructured. Accessibility: Hard to access. Governance: Very weak.
3	Low readiness	Availability: Limited. Quality: Highly inconsistent. Structure: Unstructured. Accessibility: Difficult. Governance: None.
2	Very low readiness	Availability: Barely available. Quality: Unusable without heavy effort. Structure: Raw. Accessibility: Not accessible. Governance: Absent.
1	Minimal readiness	Availability: Almost no usable data. Quality: Extremely poor. Structure: Not defined. Accessibility: Not available. Governance: None.
0	No readiness	No data exists. POC is not viable in any form.

5. Complexity

Complexity measures how difficult it is to implement the AI use case end-to-end in an enterprise environment. It is not limited to model difficulty. It evaluates both technical factors (data, model, deployment) and enterprise factors (integration, organisational change).

Weight: –10% (penalty). Some high-value Oracle use cases are naturally complex. The framework deliberately avoids over-penalising strategically important projects. The penalty is intentionally moderated.

Scoring dimensions

Dimension	What is assessed
1. Data complexity	Is data clean or messy? From a single source or multiple sources? Structured or unstructured? Multi-source, poor-quality, external data increases complexity significantly.
2. Integration complexity	Is the solution standalone, or does it require connection to ERP, SCM, HCM, and external APIs in real time?
3. Model complexity	Is the model a simple classifier, standard ML, advanced ML, deep learning, or a multi-model hybrid system?
4. Deployment effort	Is deployment a simple API call, a departmental rollout, or an enterprise-wide mission-critical workflow integration?
5. Organisational change	Does adoption require no change, minor adjustments, workflow changes in multiple teams, or major cross-functional process redesign?

Scale (0–10)

Note: Higher complexity score = harder to implement. The penalty applied in the weighted formula is –10% of this score. A score of 10 (extremely complex) subtracts 1.0 from the overall weighted score.

Score	Level	Combination
10	Extremely high	Data: Multi-source + unstructured + external + poor quality. Integration: Cross-enterprise, real-time, multiple systems. Model: Deep learning / hybrid / multi-model. Deployment: Enterprise-wide, mission-critical. Org: Major process redesign + cross-functional change.
9	Very high	Data: Multi-source, partially unstructured. Integration: Multiple enterprise systems. Model: Advanced ML. Deployment: Large-scale across business units. Org: Significant cross-team coordination.
8	High complexity	Data: Multi-source structured + some external. Integration: 2–3 systems (ERP + others). Model: Moderate–advanced ML. Deployment: Scalable across departments. Org: Process changes in multiple teams.
7	Moderately high	Data: Multiple structured datasets. Integration: Limited enterprise integration. Model: Standard ML. Deployment: Department-level. Org: Moderate workflow adjustments.
6	Medium complexity	Data: Structured + small external inputs. Integration: Light integration (one system dependency). Model: Standard ML (classification / regression). Deployment: Controlled environment. Org: Minor adjustments.
5	Moderate complexity	Data: Mostly structured, manageable quality. Integration: Minimal. Model: Simple ML. Deployment: Simple workflow integration. Org: Limited user adaptation.
4	Low–moderate	Data: Structured and clean. Integration: Single system. Model: Simple logic + light ML. Deployment: Straightforward. Org: Minimal change.

3	Low complexity	Data: Structured and ready. Integration: Minimal or none. Model: Prebuilt AI / API-based. Deployment: Plug-and-play. Org: Very low change.
2	Very low	Data: Clean and readily available. Integration: None or trivial. Model: Rule-based + minor AI. Deployment: Instant. Org: No change required.
1	Minimal complexity	Data: Fully ready, no prep. Integration: None. Model: No real ML (basic rules). Deployment: Immediate. Org: None.
0	No complexity	Not an AI problem.

6. Client Demand

Client demand measures how strongly and frequently enterprises are asking for or adopting this use case. It is not just popularity. It reflects market pull, business priority, and adoption trend.

Weight: 7.5%. Demand matters, but not all demanded use cases fit Oracle. Oracle fit takes precedence over demand. Client demand is a supporting signal, not a primary driver.

Scoring dimensions

Dimension	What is assessed
1. Market adoption level	How widely is this use case adopted or actively growing across the enterprise AI market today?
2. Frequency of client requests	How often do clients specifically ask for this use case in engagements?
3. Industry applicability	Is this relevant to only a few niche industries, specific domains, or broadly applicable across Finance, SCM, HCM, and other sectors?
4. Business urgency	Is this a nice-to-have, an important initiative, or a critical need that clients are actively investing in?

Demand bands

Client demand is assessed as one of three bands rather than a fine-grained 0–10 scale. The band maps to a numeric score for the weighted formula.

Band	Score	Combination	Interpretation	Examples
High demand	8–10	Market adoption: Widely adopted or rapidly growing. Requests: Frequent in engagements. Applicability: Cross-industry. Urgency: High — solves critical problems.	Clients actively investing. Seen as must-have or priority initiative.	Fraud detection, demand forecasting, document automation.
Medium demand	5–7	Market adoption: Moderate / growing. Requests: Occasional. Applicability: Domain-specific. Urgency: Moderate — important but not critical.	Clients are interested but not urgent. Often secondary priority.	Customer segmentation, lead scoring, predictive maintenance.
Low demand	1–4	Market adoption: Low / experimental. Requests: Rare. Applicability: Niche. Urgency: Low — nice-to-have.	Clients are not prioritising this yet. Early-stage or optional.	Advanced personalisation in niche domains, experimental AI.

7. Time to Value

Time to value measures how long it takes for an AI use case to start delivering measurable business benefits after implementation begins. It includes data preparation time, integration effort, deployment setup, and organisational adoption — not just model development time.

Weight: 7.5%. Quick wins are valued. Advisory also includes long-term transformation. This weight ensures the framework rewards speed without dismissing strategically important but slower use cases.

Note: Projects most commonly fail at the pilot-to-production transition. Short TTV reduces this risk. Long TTV increases it.

Scoring dimensions

Dimension	What is assessed
1. Data availability readiness	Is data already available and clean (accelerator) or does it require heavy preparation before any modeling can begin (blocker)?
2. Integration effort	Is the model standalone, or does it require deep integration with ERP, SCM, and other enterprise systems?
3. Deployment complexity	Is deployment a simple API call, or does it require enterprise-wide workflow redesign?
4. Organisational readiness	Does adoption require minimal process change or major change management across teams and business units?

TTV bands

Band	Score	Timeline	Combination	Examples	Risk
Short TTV	7.5–10	Weeks to 3 months	Data: Ready and clean. Integration: Minimal. Deployment: Simple API. Org: Minimal change.	Chatbot, document processing	Low
Medium TTV	5–7.4	3–6 months	Data: Available, needs prep. Integration: Moderate (1–2 systems). Deployment: Standard. Org: Some workflow change.	Attrition prediction, lead scoring	Moderate
Long TTV	0–4.9	6+ months	Data: Fragmented. Integration: Complex ERP + external. Deployment: Enterprise-wide. Org: Major process change.	Demand sensing, supply chain optimisation	High

8. Weighted Score and Prioritisation

The six dimension scores are combined into an overall weighted score using the formula below. This score drives prioritisation decisions about which use cases to take forward, which to re-scope, and which to discard.

Weight strategy and justification

Dimension	Weight	Justification
Oracle fit	35%	The engagement is Oracle Advisory, not generic AI. The use case must align with Oracle ERP/SCM/HCM and use OCI services effectively. If Oracle fit is low, the project is out of scope regardless of any other score.
Business value	25%	Advisory is ROI-driven. Clients care about revenue and cost. A use case must demonstrate high value within the Oracle ecosystem, not just in isolation.
Data readiness	15%	AI fails due to data. Even if Oracle fit is high, poor data readiness leads to POC failure. The G2 Data Readiness gate in the delivery framework is the most frequently triggered failure point.
Complexity (penalty)	−10%	Some high-value Oracle use cases are naturally complex. The framework does not over-penalise strategically important projects. The penalty is deliberately moderated to −10% rather than a steeper discount.
Client demand	7.5%	Demand matters but not all demanded use cases fit Oracle. Oracle fit takes precedence over demand. Client demand is a supporting signal that confirms market viability.
Time to value	7.5%	Quick wins are valued. Advisory also includes long-term transformation programmes. This weight rewards speed without dismissing strategically important but longer-horizon use cases.

Formula

Weighted score = (Oracle fit × 0.35) + (Business value × 0.25) + (Data readiness × 0.15) − (Complexity × 0.10) + (Client demand × 0.075) + (Time to value × 0.075)

Score interpretation

Score range	Recommendation	Action
8.0 – 10.0	Strong candidate	Proceed to detailed feasibility and POC gate assessment. Prioritise in pipeline.
6.5 – 7.9	Good candidate	Viable for POC with minor adjustments. Address the weakest dimension before committing full resource.
5.0 – 6.4	Moderate — address gaps	Identify and resolve the weakest dimension before proceeding. Consider adjacent use cases with better scores.
< 5.0	High risk — review	Do not proceed to POC. Fundamental blockers exist in one or more dimensions. Re-scope or replace the use case.

Automatic blockers

Regardless of overall score, the following dimension scores act as automatic blockers. A use case scoring below these thresholds cannot enter the POC delivery framework.

Dimension	Blocker threshold	Consequence
-----------	-------------------	-------------

Oracle fit	< 5.0	Use case is out of scope for Oracle Advisory. Re-scope to better leverage Oracle systems or replace with a different use case.
Data readiness	< 5.0	POC cannot begin. Data infrastructure must be addressed first. The G2 gate in Process Map A will fail. Raises overall delivery risk significantly.

What high-scoring and low-scoring use cases look like

	High-ranking use cases	Low-ranking use cases
Oracle integration	Strong ERP/SCM/HCM integration with native OCI pipeline	Standalone AI with no Oracle integration
Data source	Data available in Oracle ERP or SCM	External data-heavy systems outside Oracle
Business impact	High financial impact on core Oracle workflows	Limited impact on non-Oracle processes
Examples	Invoice anomaly detection, inventory optimisation, demand forecasting	Generic chatbots, external data analytics, niche standalone AI

9. Alignment with POC Delivery Framework

The scoring dimensions map directly to specific decision gates in the AI POC Delivery Framework (Process Map A). Low dimension scores are early warnings that a particular gate is likely to fail during delivery.

Score dimension	POC gate	Risk if score is low
Oracle fit	G3 — Oracle fit gate	G3 will fail. Re-scoping or platform change required before modeling investment is committed.
Data readiness	G2 — Data readiness gate	G2 will fail. Enrichment loop triggered. Modeling cannot begin until data is resolved. Most common failure point.
Business value	G1 — Business value gate	G1 will fail. Value hypothesis must be strengthened before any technical delivery begins.
Complexity	G4 — Complexity vs TTV gate	G4 may fail if effort-to-value ratio is poor. MVP scope should be defined to reduce complexity before committing to full delivery.
Org readiness (part of complexity)	G5 — Org readiness gate	G5 will fail if organisational readiness is low. Stakeholder alignment and change management plan must be in place.
Overall score	G7 — Scalability gate	A low overall score signals the use case is unlikely to survive the full G1–G7 gate sequence. Re-evaluate before POC entry.

10. Important Notes

- Financial estimates and score thresholds are indicative starting points. Detailed calibration is required before the framework is finalized for client delivery.
- Scores should be reviewed and updated at each stage gate as new information becomes available during the POC lifecycle. The scoring document is a living input to the gate decisions.
- The framework is designed for Oracle-aligned engagements. Use cases evaluated outside an Oracle advisory context should use a modified weighting model that removes the Oracle-first bias.
- Complexity is scored as a penalty. A high complexity score (hard to implement) lowers the overall score, incentivising realistic scope definition and MVP thinking over over-ambitious delivery plans.